

[GitHub Repository](#)

Sports statistics have always been a passion of mine, whether it be analyzing quarterback ratings, advanced individual baseball statistics, or measuring the odds of certain events occurring. However, the sport that is nearest and dearest to my heart is college basketball, because of the rivalry and unpredictability of the sport. The dataset I chose for this set detailed the statistics on all NCAA Division 1 basketball teams from the 2012-13 season until the 2022-23 season. (Note that the 2019-20 season was not included with this due to no postseason play.) I wanted to create a model that would classify teams as having the potential to “go deep,” or win their first three games, in the NCAA tournament. For reference, the NCAA Tournament is a single elimination tournament that takes 68 teams from across the country, pits them against each other at neutral sites, and over the course of three weekends, a champion is crowned. The ideal model would predict eight teams correctly every year, as only eight teams win at least three games every year in this tournament. To do this, I created a logistic regression model using stratified splits and considered F1 scores of each split to determine the best models.

To begin, the dataset contained 24 unique features and 3523 unique observations. Both numerical and categorical features were present. The majority of numerical variables had to do with average gameplay statistics, with there being specific features for the number of games played, the number of games won, the number of teams they won against that were above the “bubble”, the year their season took place, and their NCAA tournament “seeding” position if they had qualified. The categorical variables included the names of the teams, their current conference membership, and their postseason finish if they had qualified.

```
[4] cbb.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3523 entries, 0 to 3522
Data columns (total 24 columns):
 #   Column          Non-Null Count  Dtype  
---  --
 0   TEAM            3523 non-null   object  
 1   CONF            3523 non-null   object  
 2   G               3523 non-null   int64   
 3   W               3523 non-null   int64   
 4   ADJOE           3523 non-null   float64  
 5   ADJDE           3523 non-null   float64  
 6   BARTHAG         3523 non-null   float64  
 7   EFG_O           3523 non-null   float64  
 8   EFG_D           3523 non-null   float64  
 9   TOR             3523 non-null   float64  
10  TORD            3523 non-null   float64  
11  ORB             3523 non-null   float64  
12  DRB             3523 non-null   float64  
13  FTR             3523 non-null   float64  
14  FTRD            3523 non-null   float64  
15  2P_O            3523 non-null   float64  
16  2P_D            3523 non-null   float64  
17  3P_O            3523 non-null   float64  
18  3P_D            3523 non-null   float64  
19  ADJ_T           3523 non-null   float64  
20  WAB             3523 non-null   float64  
21  POSTSEASON      680 non-null    object  
22  SEED            680 non-null    float64  
23  YEAR            3523 non-null   int64   
dtypes: float64(18), int64(3), object(3)
memory usage: 660.7+ KB
```

Figure 1: All features found in cbb.csv

Cleaning my dataset took longer than expected, primarily due to the fact I wasn't sure what to do with the large number of observations I had. It is important to note that out of all observations, only 680 of them were teams that qualified to play in the NCAA tournament. At first, I wanted to keep them in because I wanted to try and consider all teams, but this skewed the dataset so much that only two percent of teams met the criteria for "going deep." As a result, I adjusted my dataset to contain only teams that qualified for the NCAA tournament. I also went on to create two new variables, "DEEP" and "P7", from existing features present in the dataset. "DEEP", or our target variable, was a Boolean feature that indicated whether or not a team made it to the Elite 8 or further in the postseason. "P7" was a Boolean feature as well that indicated whether or not a team was from a "Power" conference or not. These conferences tend to have more funding and resources than your average Division 1 school, so they typically are more likely to make the postseason and to perform well in the postseason. To make both of these features, I utilized the .isin() function alongside the feature I was drawing my information from to extract the data I needed. After eliminating a few features I found to be redundant and converting all features to floats, I was left with 19 data features and my target feature.

```
# Extract all observations where 'SEED' is not nan
cbb = cbb.dropna(subset=['SEED'])

cbb['CONF'].unique()
# Power 7 Conferences are ACC, B10, B12, SEC, BE, P12, and Amer
# Wish to create a column indicating if an entry belongs to P7 or not
cbb['P7'] = cbb['CONF'].isin(['ACC', 'B10', 'B12', 'SEC', 'BE', 'P12', 'Amer'])

# My target variable for this project is going to be teams "going deep" in the
# NCAA tournament. To go deep, you need to win at least 3 games, or make it to
# the Elite 8, denoted by E8 in the POSTSEASON feature set.
# I am going to create a feature named DEEP that meets this category.
print(cbb['POSTSEASON'].unique())

# Teams denoted as Champion, 2ND, F4, or E8 meet this category.
cbb['DEEP'] = cbb['POSTSEASON'].isin(['Champions', '2ND', 'F4', 'E8'])
```

Figure 2: Cleaning methods used to tidy up and expand the dataset.

```
# Removing TEAM feature, as supposed to view these as blind resumes
cbb = cbb.drop(['TEAM', 'CONF', 'YEAR', 'POSTSEASON', 'G'], axis=1)

cbb.info()

<class 'pandas.core.frame.DataFrame'>
Index: 680 entries, 0 to 3227
Data columns (total 21 columns):
#   Column      Non-Null Count  Dtype
---  -
0    W           680 non-null    int64
1    ADJOE       680 non-null    float64
2    ADJDE       680 non-null    float64
3    BARTHAG     680 non-null    float64
4    EFG_O       680 non-null    float64
5    EFG_D       680 non-null    float64
6    TOR         680 non-null    float64
7    TORD        680 non-null    float64
8    ORB         680 non-null    float64
9    DRB         680 non-null    float64
10   FTR         680 non-null    float64
11   FTRD        680 non-null    float64
12   2P_O        680 non-null    float64
13   2P_D        680 non-null    float64
14   3P_O        680 non-null    float64
15   3P_D        680 non-null    float64
16   ADJ_T       680 non-null    float64
17   WAB         680 non-null    float64
18   SEED        680 non-null    float64
19   P7          680 non-null    bool
20   DEEP        680 non-null    bool
dtypes: bool(2), float64(18), int64(1)
memory usage: 107.6 KB
```

Figure 3: Final dataset contents upon EDA completion.

As I had only two options that could be present for my feature, I sought to treat this as a classification problem. I wanted to exploit thresholds in the features that the majority of qualifying teams would exceed. Before modeling the dataset, I tried to find which features were highly correlated with the target feature. No features had an extremely high correlation, so I had to lower the correlation threshold down to 0.25 to procure more than five features. The six features I had left were the number of wins the team earned, their adjusted offensive efficiency per 100 possessions, their defensive efficiency per 100 possessions, their BARTHAG (which represents the percentage they could win against an average team on a neutral court), the number of wins above bubble they earned, and their tournament seeding.

```
[7] corr = cbb.corr(method='pearson')

    high_corr_features = corr.index[abs(corr['DEEP']) > 0.25]
    print(high_corr_features)

Index(['W', 'ADJOE', 'ADJDE', 'BARTHAG', 'WAB', 'SEED', 'DEEP'], dtype='object')
```

Figure 4: Features that met the 0.25 threshold to be considered in the model.

I wanted at first to use a decision tree setup to model whether these tournament teams could go deep or not. However, as I began to try and apply the model, I discovered this was not a dataset I could apply this onto. The main reason for this was that the NCAA Tournament is, by nature, incredibly unpredictable. Many high seeded teams lose games early on, eliminating them from the tournament. That is what makes this tournament so interesting to the masses, the prospect a team from Nowhereville, Statescatchewan, can knock off a big-name power conference team like Kansas, Kentucky, or Duke. As a result, I chose not to document this attempt and refocused on creating a regression model. I chose to use logistic regression for this model, as teams that performed better in the tournament usually had their results clustered closer together than the ones that didn't.

It is important to note that with this dataset, even after I removed teams that did not qualify for the tournament, there was still a high level of weight shifted towards teams that didn't go deep, with 88.2% of teams not meeting this criterion. As a result, for this model to work, I had to use stratified splits to evenly distribute the teams into the train and test sets used for modeling. To find the results of interest, I considered three criteria to see which best result of these three produced the most accurate results. Within each k-split for loop, I chose to examine the results of each model produced and store the best one. I was then able to retrain the stored model as my final model. To see this, please refer to code block 10 in `azinkSelectionAndTraining.ipynb`, which can be found in my GitHub repository for this project.

To test the model's performance, I chose to use another dataset from within the original files. This set was from the previous season (2023-24), but was incomplete, as it was published before the NCAA Tournament was completed, therefore having no target variable. Since the results of this tournament are in the public eye now, I chose to use this set as my evaluation set, rather than generate random observations.

Tree selection was also a struggle as well, as I believed that this would have been a more optimal way to select these teams. I think a decision tree application would be better suited for a playoff structure like what the NBA uses, where teams play against each other over the course of a series of games, with a team having to win 3 or 4 games against their opponent in order to advance. Finally, I would like to think that separating teams that belong to Power conferences and non-power conferences from each other would allow for more accurate predictions. Non-power conference teams typically do not perform as well, but there are a few outliers that lead me to believe they should still be considered in the analysis of this dataset. This could be done by using the P7 feature created in the EDA section of the project.